

ACI-CO Repository Architecture

Colombian Actuarial Climate Index (ACI-CO) pipeline

August 6, 2026

Abstract

This document is a LaTeX companion to `ARCHITECTURE.md` and `CLAUDE.md` at the repository root. It describes how the ACI-CO codebase is organized: the scientific target it computes, the actual data flow between scripts (as opposed to the README’s idealized version), where the repository’s documentation is authoritative versus stale, and how to read already-processed outputs from a notebook (e.g. Google Colab) using the `aci_lib` data-access library introduced alongside this document. It is a description of the codebase as it stands, not a specification – where the code and the docs disagree, the code wins.

Contents

1	What this repository is	2
2	The scientific target: ACI-CO	2
2.1	Extensions beyond the base ACI framework	2
2.2	Regions actually computed	3
3	Repository layout	3
4	Actual data flow	4
4.1	Two parallel “correctness check” implementations	5
4.2	Forecasting: two families plus shared modules	5
5	Documentation status	5
6	Discrepancies between README/docstrings and reality	6
7	The git tracking gap	6
7.1	Known secrets in the repository	7
8	The <code>aci_lib</code> data-access library	7
8.1	Why it exists	7
8.2	Important caveat: data is not in git	7
8.3	Quickstart in Google Colab	8
8.4	Scope	8
9	Cleanup priorities (standing list)	8

1 What this repository is

This is a research pipeline that builds the **Colombian Actuarial Climate Index (ACI-CO)** from ERA5(-Land) reanalysis data (1961–2024): a Colombian adaptation of the Actuaries Climate Index (ACI) framework, extended with ENSO calibration, a bootstrap uncertainty band, multi-dataset robustness checks, and validation against Colombia’s national disaster registry (UNGRD).

It is an active research codebase, not a packaged library: there is no test suite, several parallel/experimental implementations of the same pipeline stage coexist deliberately, and a non-trivial share of the *working* pipeline was never committed to git (Section 7).

Authoritative methodology source. The methodology (how T90, T10, Rx5day, CDD and WP are computed; how the composite, ENSO decomposition, and bootstrap bands work) is described in `articles/aci_co_submission/article1.tex`, §3–§7 – *not* the older root-level `articles/article1.tex` (a superseded LOWESS-based draft), and not the empty stubs in `docs/methodology.md` / `docs/results.md` / `docs/zones_description.md`.

2 The scientific target: ACI-CO

The ACI-CO composite is the unweighted mean of five signed, standardized components (divided by $K = 5$; an earlier draft divided by 6 for a placeholder sea-level component, since dropped for lack of tide-gauge coverage):

Component	Meaning	Sign	Threshold
T90	Fraction of days/month with temperature above the 1961–1990 baseline 90th percentile	+	Empirical P90 per grid cell × calendar month
T10	Fraction of days/month with temperature below the baseline 10th percentile	–	Empirical P10
Rx5day	Max 5-day precipitation total per month, z -scored vs. 1961–1990 climatology	+	z -score vs. monthly climatology
CDD	Annual max consecutive dry days (<1 mm), interpolated onto the monthly axis	+	—
WP	Wind-power ($\propto U^3$) exceedance	+	Empirical P90 (current); replaces an earlier parametric $\mu + 1.28\sigma$ threshold, discarded because WP is right-skewed

2.1 Extensions beyond the base ACI framework

- **ENSO-neutral decomposition** (§4): asymmetric ARIMAX regression per component/region on Niño-3.4 ($E^+ = \max(E, 0)$, $E^- = \min(E, 0)$) plus a linear trend and $\text{AR}(p)$ errors, separating secular warming from ENSO-driven variability.

- **Bootstrap uncertainty bands** (§4.4): block-resample ($L = 12$ mo) the 1961–1990 baseline composite into $B = 500$ (national) / $B = 200$ (regional) null-model replicates, testing whether the observed trend exceeds baseline-period variability. Replaced an earlier LOWESS residual-envelope approach.
- **Multi-dataset robustness** (§5): ERA5-Land vs. ERA5 standard vs. CHIRPS.
- **UNGRD validation** (§7): negative-binomial regression of disaster counts (floods, landslides, windthrow, vegetation fires – 49,734 events, 1998–2022) on ACI-CO components.

2.2 Regions actually computed

Despite the paper conceptually discussing 5 natural regions and 32 departments, the implementation aggregates **4 units**: national (area-weighted, all of Colombia) and 3 departments chosen for insurance relevance – **Antioquia**, **Cundinamarca-Bogotá**, and **Valle del Cauca**. Shapefiles also exist for Bogotá, Medellín, Cali, San Andrés/Providencia, Pacífico, and Amazonas, but the corresponding `data/processed/anomalias_{bogota,medellin,san_andres_providencia}/` directories are empty – never populated.

3 Repository layout

Path	Contents
<code>src/scripts/</code>	The original OPT pipeline: download, percentile, and anomaly-calculation scripts (<code>calcular_*</code> , <code>ecmwf_descarga.py</code> , <code>unir_archivos.py</code> , <code>graficas.py</code> , ...). Spanish-language docstrings dominate here.
<code>src/forecast_scripts/</code>	Two independent forecasting families (AutoTS ensembles vs. SARIMAX/ETS + ENSO) plus <code>common/</code> , the shared modules extracted from both (Section 4.2). English-language, newer code.
<code>src/dashboard/</code>	<code>app.py</code> – a read-only Streamlit dashboard that reads Stage 3/5 CSVs directly; not a pipeline stage.
<code>src/utils/</code>	Shared helpers (<code>get_available_years</code> , <code>get_cached_shapefile</code>), consolidated from near-identical copies previously duplicated across the <code>calcular_*</code> scripts.
<code>src/aci_lib/</code>	Read-only data-access library for notebooks (Section 8) – added alongside this document, not part of the original OPT/forecast pipeline.
<code>data/raw/</code> , <code>data/processed/</code> <code>articles/</code>	Pipeline inputs/outputs, entirely gitignored except for <code>.gitkeep</code> placeholders (Section 7). Three papers at different stages: the current submission (<code>aci_co_submission/</code>), a superseded LOWESS draft (<code>article1.tex</code>), and a separate forecasting report (<code>AnnualICAForecast_Report.tex</code>).
<code>docs/</code>	Mostly empty stubs (<code>methodology.md</code> , <code>results.md</code> , <code>zones_description.md</code>); the one substantive file is <code>Diccionario.md</code> , a ~2400-line data-source vetting document.
<code>tests/</code>	Placeholder only (<code>test.txt</code>); there is no test suite or CI.

4 Actual data flow

The README describes an idealized 8-step flow; several of the scripts it names do not exist (Section 6). The flow below is reconstructed from what the scripts actually read and write.

Stage	Script(s)	Output
0. Download	<code>ecmwf_descarga.py</code> (canonical), <code>sst_descarga.py</code> , <code>get_aux_data.py</code> , IDEAM station data via <code>pyreadr</code>	<code>data/raw/era5/*.grib</code> , SST GRIB, <code>data/raw/auxiliary/</code> , <code>data/raw/ideam/*.rds</code>
1. Merge / resample	<code>unir_archivos.py</code> ; independently, the untracked <code>run_aci_colombia.py</code> re-merges raw GRIB into a from-scratch reimplementaton	<code>era5_daily_combined_{tmp,rain, wind}.nc</code> ; <code>data/processed/aci_daily/</code> , <code>aci_colombia_1961_2024.csv</code>
2. Percentiles (1961–1990 baseline)	<code>calcular_percentil_- {temperatura,lluvia, viento}.py</code>	<code>era5*_percentil.nc</code> , <code>percentiles*.csv</code>
3. Anomalies, per region	<code>calcular_anomalias_- {temperatura,lluvia, viento}.py</code> , orchestrated by <code>calcular_anomalias_- regiones.py</code> over <code>data/shapefiles/*.shp</code>	<code>data/processed/anomalias_- <region>/anomalias_{temperature, precipitation,drought, wind}_combined.csv</code> (populated: antioquia, cali, colombia, cundinamarca_bogota, valle_cauca; empty: bogota, medellin, san_andres_providencia)
4. Daily variants	<code>produce_daily_- anomalies.py</code> , <code>append_2025_2026.py</code> (untracked)	<code>data/processed/daily_- anomalies/*.xlsx</code>
5. ENSO / SST index	<code>anomalias_sst_daily.py</code> , <code>compareenso_flags.py</code>	<code>sst_index_colombia_pacific.csv</code> , <code>enso_flag_comparison.csv</code> , validated against NOAA CPC <code>oni.ascii.txt</code>
6. Downscaling	<code>downscale_era5_- temperature.py</code> , <code>downscale_era5_tp.py</code> (RandomForest, DEM/NDVI/IDEAM)	<code>data/processed/downscaled/</code> , validated in <code>data/processed/era5_- ideam_comparison/</code>
7. Visualization / export	<code>graficas.py</code> , <code>generate_- correlation_matrices.py</code>	<code>articles/graficas/<region>/*.png</code> , xlsx exports, correlation heatmaps
8. Forecasting	Two independent families – see Section 4.2	<code>articles/graficas/forecast_*/</code>

`data/indices/` and `data/zones/` exist but are effectively empty – no script currently writes to those exact paths.

4.1 Two parallel “correctness check” implementations

There are effectively two independent implementations of the core ACI algorithm in this repository:

- **The OPT pipeline** (`src/scripts/calcular_*`, tracked in git, the README’s documented flow) – the original project pipeline, per-region, percentile/anomaly based.
- `run_aci_colombia.py` (untracked) – a from-scratch reimplementaion of the ACI-Python reference algorithm directly on raw GRIB, used to cross-validate the OPT pipeline’s output. `compare_aci_colombia.py` and `compare_repos.py` (also untracked) diff the two and document the algorithmic differences.

Neither is “the old version” of the other – they are independent implementations kept deliberately in sync for cross-validation, and this comparison grounds the paper’s “validation against x-ACT” discussion (§8).

4.2 Forecasting: two families plus shared modules

Two independent forecasting families exist side by side and are *not* a superset/subset of each other:

- `forecast_ica_{monthly,daily}.py` – AutoTS ensembles.
- The ETS(X)_* family – statsmodels SARIMAX/ETS with an ENSO exogenous regressor.

`src/forecast_scripts/common/` holds logic that was byte-identical or near-identical across these scripts before consolidation:

Module	Extracted from	Used by
<code>data_loading.py</code>	Byte-identical in <code>forecast_ica_monthly.py</code> and <code>ETS_ica_forecast.py</code>	Both, via thin same-named wrapper methods
<code>climate_index.py</code>	Near-identical in all three monthly forecasters	All three
<code>stationarity.py</code>	Near-identical in <code>ETS_ica_forecast.py</code> / <code>ET SX_ica_forecast.py</code>	Both
<code>diagnostics.py</code> (ModelError-Diagnostics)	Moved unchanged from the old <code>error_diagnostics.py</code>	Not currently called by anything

Deliberately *not* merged (structurally parallel, not literal duplicates): the `ONIDataHandler` / `ONIDataHandlerDaily` ENSO-fetching classes, the SARIMAX+Lasso fitting methods, the AutoTS setup (same shape, different hyperparameters), and all `visualize_forecasts`-style plotting methods (4 separate implementations, each with different overlays).

5 Documentation status

Document	Status
docs/methodology.md, docs/results.md, docs/zones_- description.md	Empty stubs. Methodology lives in articles/aci_co_submission/article1.tex instead.
docs/explicacion_- excel.md	One truncated sentence, no body.
docs/data_dictionary.md	Early, incomplete fragment (HUMBOLDT-only) of docs/Diccionario.md.
docs/Diccionario.md	The real, complete data-source vetting document (~2400 lines, 20 candidate sources evaluated). Canonical over <code>data_dictionary.md</code> .
README.md	Describes an idealized 8-step flow; two referenced scripts do not exist.

6 Discrepancies between README/docstrings and reality

- README step 1 says to run `descargar_datos.py` – it does not exist; `ecmwf_descarga.py` is the actual entry point.
- README step 6 says to run `sealevel.py` – does not exist. `sealevel2.py` exists but does not do sea-level analysis; it is a Mann-Kendall trend + Q-Q normality check on precipitation-like series with leftover Jupyter cell markers, suggesting it was pasted from a notebook and misnamed. Sea level is not part of the current ACI-CO methodology.
- `src/scripts/earthhub_descarga.py` has a whole-file indentation error and cannot run as-is, on top of embedding a live-looking credential (Section 7.1).
- `src/scripts/display.py` is not a plotting module – it is a Greece-precipitation download script (also embeds the same credential) that produced a 562 MB `.nc` file at the repo root as a side effect.
- Stale internal docstrings/comments: `era5datasets.py` says `""""compare_land_only.py""""`; `get_aux_data.py`'s header says `get_auxiliary_data.py`. Trust the filename/git path.
- `graficas_altaresolucion.py` is a 4-line dead stub.
- `ETSX_seasonal_motif_forecast.py` (the “broken two-stage Lasso→ETS” approach `ETSX_daily_ica_forecast.py`'s docstring says it replaced) has been deleted – confirmed unused before removal.
- `example_grid_search_era5.py` imports `ONIDataHandler` from `ETSX_ica_forecast.py` expecting methods that do not exist on that class – a pre-existing break, unrelated to the `forecast_scripts` modularization.
- `articles/AnnualICAForecast_Report.tex` references `scripts/daily_forecast_to_annual.py` and `scripts/ar1_vs_ets_final.py` – neither exists anywhere in the repository. Treat that report's pipeline as documented-but-unimplemented.

7 The git tracking gap

A significant portion of the *current, more-correct* pipeline exists only in the working tree and has never been git added: `run_aci_colombia.py`, `compare_aci_colombia.py`, `compare_-`

`repos.py`, `create_region_shapefiles.py`, `probe_2025.py`, `produce_daily_anomalies.py`, `append_-2025_2026.py`, `reproduce.py`, `plot_repo_anomalies.py`, and most of `articles/` (the `aci_co_-submission/` folder, `produce_paper_plots.py`, `build_validation_outputs.py`, etc.).

Conversely, things that *are* committed include generated junk (`presentation.aux/log/out`, `texput.log`, `output.log`, `comparison.png`, `storm_daniel.png`, `.vs/`, `proyecto-indices/`, root-level `.nc/.csv` data files) and a `.gitignore` that was itself accidentally overwritten with a PowerShell heredoc command instead of its output.

Untracked here often means “uncommitted work in progress,” not “disposable.” Do not delete untracked files without checking which side of this split they are on.

7.1 Known secrets in the repository

Two Earth Data Hub PAT tokens are hardcoded and **committed** in `src/scripts/display.py` (line 13) and `src/scripts/earthhub_descarga.py` (line 15) – the same token, duplicated. A CDS API key also appears in `GRID_SEARCH_AND_ERA5_GUIDE.md` (line 44). These require rotation/history-scrubbing independent of any code change. New download scripts should follow the pattern in `ecmwf_descarga.py` / `sst_descarga.py`: `cdsapi.Client()` reads credentials from the user’s local `~/.cdsapirc`, nothing embedded in source.

8 The `aci_lib` data-access library

`src/aci_lib/` is a small, dependency-light Python package that turns `data/processed/` into a browsable, loadable catalog, so that already-computed pipeline outputs can be pulled into a notebook (locally or in Google Colab) without knowing the directory’s internal layout by heart.

8.1 Why it exists

`data/processed/` mixes single tabular files, single gridded NetCDF files, and long runs of per-year or per-year-month NetCDF files that really form one logical time series (e.g. 768 files under `anomalias_-colombia/anomalies_temperature_1961_1.nc` through `_2024_12.nc`). `aci_lib` groups the latter automatically and exposes a uniform three-function API:

- `list_datasets()` – a `DataFrame` of every dataset name, kind, file count, and (where known) a description sourced from this document’s data-flow tables.
- `describe(name)` – the underlying file list and kind for one dataset.
- `load(name)` – returns a `pandas.DataFrame` for CSV/XLSX datasets, an `xarray.Dataset` for single or multi-file NetCDF datasets (opened lazily via `xarray.open_mfdataset`), or a file path for images/raw text files.

8.2 Important caveat: data is not in git

`data/processed/` is entirely gitignored (Section 7) – only a `.gitkeep` placeholder is tracked. A plain git clone of this repository, including in Colab, brings the *code* in `src/aci_lib/` but *not* the data it reads. The data must be supplied separately: mount it from Google Drive, upload and unzip an archive of `data/processed/`, or copy it in via any other means available in the notebook environment.

8.3 Quickstart in Google Colab

```
# 1. Get the code (data/processed/ is gitignored, so this clone
# does NOT include the data -- only src/aci_lib/).
!git clone https://github.com/<org>/aca_indice_climatico_opt-main.git
import sys
sys.path.insert(0, "/content/aca_indice_climatico_opt-main/src")

# 2. Supply the data. Pick ONE of:
# (a) mount Drive, if data/processed/ was uploaded there beforehand:
from google.colab import drive
drive.mount("/content/drive")
DATA_DIR = "/content/drive/MyDrive/aca_data/processed"

# (b) upload + unzip a local copy:
# from google.colab import files
# files.upload() # e.g. processed.zip
# !unzip -q processed.zip -d /content/aci_data
# DATA_DIR = "/content/aci_data/processed"

import aci_lib
aci_lib.set_data_dir(DATA_DIR)

# 3. Browse and load.
catalog = aci_lib.list_datasets()
catalog.head(20)

national_composite = aci_lib.load(
    "anomalias_colombia/anomalies_temperature_combined"
) # -> pandas.DataFrame

gridded_series = aci_lib.load(
    "anomalias_colombia/anomalies_temperature"
) # -> xarray.Dataset, lazily concatenated from ~768 monthly files
```

Run locally, inside a checkout of this repository, `set_data_dir()` is unnecessary: `aci_lib.get_data_dir()` walks up from `src/aci_lib/` to find the repository root (marked by `CLAUDE.md`) and locates `data/processed/` on its own.

8.4 Scope

`aci_lib` currently covers everything under `data/processed/` only – it does not reach into `articles/*/graficas/paper_figures/` (the paper’s final figures/tables) or `data/raw/`. Both are reasonable follow-up extensions if a notebook workflow needs them, but were kept out of scope to keep the catalog’s grouping logic simple and its descriptions trustworthy.

9 Cleanup priorities (standing list)

Ordered roughly by urgency/impact; see `ARCHITECTURE.md` for the canonical, living version of this list.

1. **Secrets:** rotate the Earth Data Hub PAT tokens and the CDS key (Section 7.1); scrub from git history if the two committed copies need to be fully removed.
2. **Commit the real pipeline:** the untracked files listed in Section 7.

3. **Purge committed generated junk:** build artifacts and data files committed at the repo root instead of under `data/`.
4. **Fix `.gitignore`:** currently wraps the real ignore rules in a literal PowerShell heredoc command that got committed instead of its output.
5. **Delete disk-only bloat** (not committed, safe to remove locally): environment archives, unrelated test data, `__pycache__/`.
6. **Consolidate duplicated helpers** into `src/utils/` (mostly done – see `get_available_years`, `get_cached_shapefile`).
7. **Resolve naming/content mismatches:** `display.py`, `sealevel2.py`, stale internal docstrings.
8. **Decide the fate of superseded scripts:** `graficas_altaresolucion.py`, `docs/data_dictionary.md`, root `articles/article1.tex`.
9. **Fill or drop the empty doc stubs** in `docs/`.
10. **Populate or drop** the empty `data/processed/anomalias_{bogota,medellin,san_andres_providencia}/` directories and the unused `data/indices/`, `data/zones/` locations.

Confirm scope with a maintainer before acting on any of these – several things that look like dead scratch code are load-bearing work that simply never got committed (Section 7).